

ProjectDb

A Key Value Storage Library

Lieyang Chen (lc3548)
Mengwen Li (ml4643)
Shengtian Mao (sm4954)

<https://github.com/mli9502/ProjectDb>

Acknowledgements and References

Martin Kleppmann: Designing Data-Intensive Applications

LevelDB: <https://github.com/google/leveldb>

Trending YouTube Video Statistics:


<https://www.kaggle.com/datasnaek/youtube-new?select=USvideos.csv>

Introduction

- In this project, we implemented a key-value storage engine library
- Provides a better performance comparing to directly writing file/reading file from disk.
- Provide a user friendly interface.
- A key-value storage engine could be used as a base for distributed NoSQL databases.

Algorithm

LSM-Tree (Log-structured merge-tree)

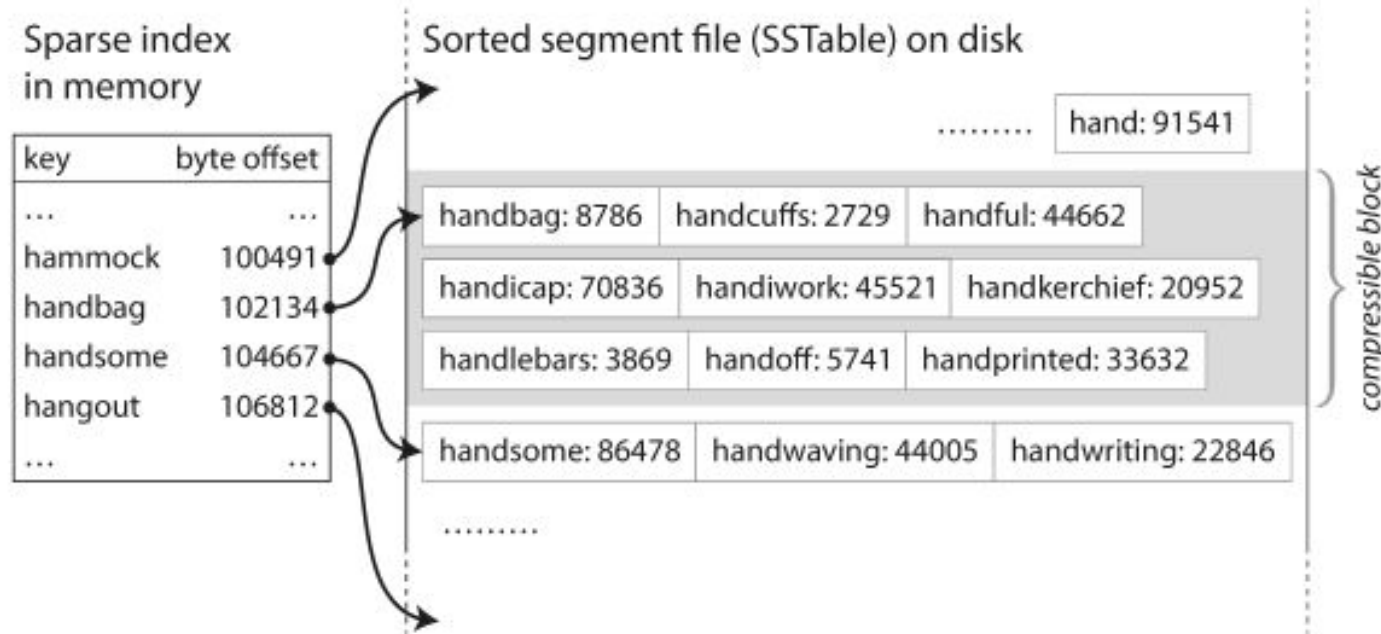


source: https://en.wikipedia.org/wiki/Log-structured_merge-tree#/media/File:LSM_Tree.png

Basic Terminologies

- **MemTable**: A table containing sorted key-value entries stored in memory.
- **SSTable**(aka segment file): A table containing sorted key-value entries stored in disk. (Created by flush Memtable into disk).
- **Segment**: A small-size block of sorted key-value entries stored in disk.
Memory reads in a segment instead of entire SSTable when performing read operations. (SSTable is composed of many segments).
- Sparse Index Table (**SSTableIndex**): A table that keeps track of the position of the beginning of each Segment.

An intuitive Graph



source: Designing Data-Intensive Applications: p76 - p79

Workflow

- When a write comes in, add it to an in-memory table sorted by key.
(MemTable)
- When MemTable gets bigger than some threshold, write it out to disk as an **SSTable** file.
- In order to serve a read request, first try to find the key in MemTable, then in the most recent on-disk **SSTable**, then in the next-older **SSTable**, etc.
- From time to time, run a **merging and compaction** process in the background to combine segment files and to discard overwritten or deleted values.

APIs

- `void set(const string& key, const string& value);`
- `string get(const string& key);`
- `string remove(const string& key);`

Serialization/Deserialization


```
71  /**
72   * Serializable follows the following recursive definitions:
73   *
74   * SerializableBase = Trivial |
75   * SerializableBase = SerializableUserDefinedType
76   * Serializable = SerializableBase |
77   * Serializable = Pair<Serializable, Serializable> |
78   * Serializable = Container<Serializable>
79  */
80  // Define type trait for the basic serializable unit.
81  template <typename T>
82  struct serializable_base_trait : std::false_type {};
83
84  template <SerializableBase T>
85  struct serializable_base_trait<T> : std::true_type {};
86  // Define type trait for generic serializable.
87  template <typename T>
88  struct serializable_trait : serializable_base_trait<T> {};
89
90  template <Pair T>
91  struct serializable_trait<T>
92  : conjunction<serializable_trait<remove_const_t<typename T::first_type>>,
93               serializable_trait<typename T::second_type>> {};
94
95  template <Container T>
96  struct serializable_trait<T> : serializable_trait<typename T::value_type> {};
97
98  template <typename T>
99  concept Serializable = serializable_trait<T>::value;
```

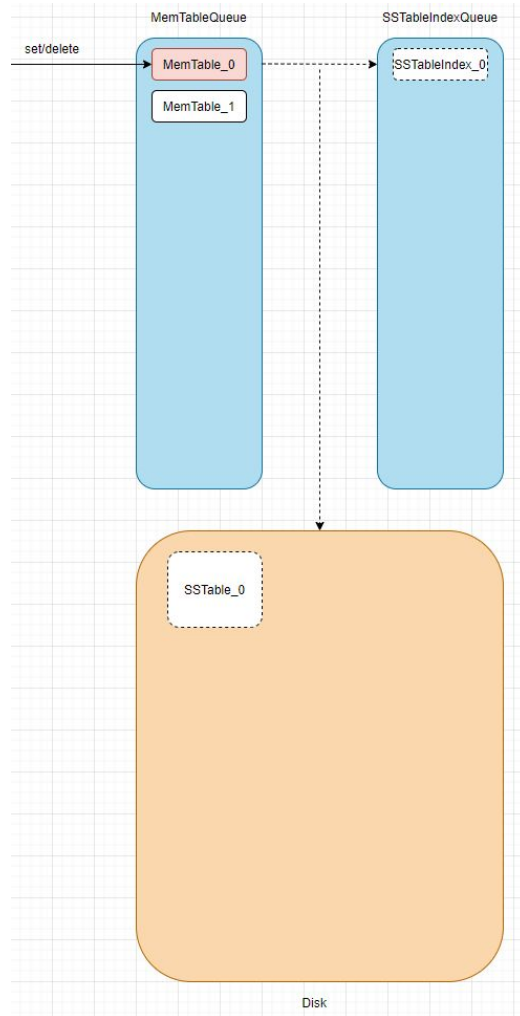
Serialization/Deserialization


```
136     template <SerializableUserDefinedType T>
137     class SerializationWrapper<T> : public impl::SerializationWrapperBase<T> {
138     public:
139         explicit SerializationWrapper(const T& t)
140             : impl::SerializationWrapperBase<T>(t){};
141         explicit SerializationWrapper(T&& t)
142             : impl::SerializationWrapperBase<T>(move(t)) {}
143
144         void operator()(ostream& os) && {
145             const T& t = this->getCRefT();
146             t.serializeImpl(os);
147             if (!os) {
148                 log::errorAndThrow("Failed to serialize Serializable data!");
149             }
150         }
151     };
```


Serialization/Deserialization

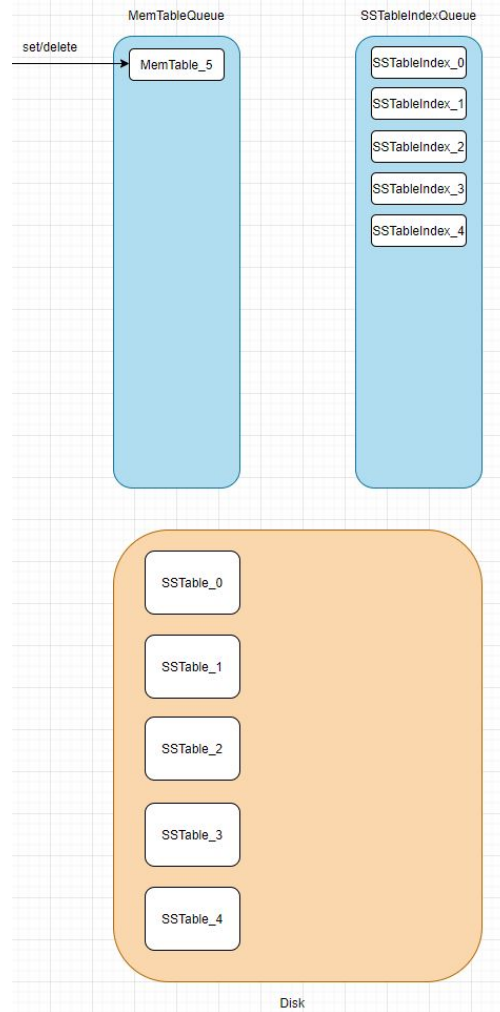
```
169     template <SerializablePair T>
170     class SerializationWrapper<T> : public impl::SerializationWrapperBase<T> {
171     public:
172         explicit SerializationWrapper(const T& t)
173             : impl::SerializationWrapperBase<T>(t){};
174         explicit SerializationWrapper(T&& t)
175             : impl::SerializationWrapperBase<T>(move(t)) {}
176
177         void operator()(ostream& os) && {
178             const T& t = this->getCRefT();
179             SerializationWrapper<
180                 typename remove_const<typename T::first_type>::type>(t.first)(os);
181             SerializationWrapper<typename T::second_type>(t.second)(os);
182         }
183     };
```

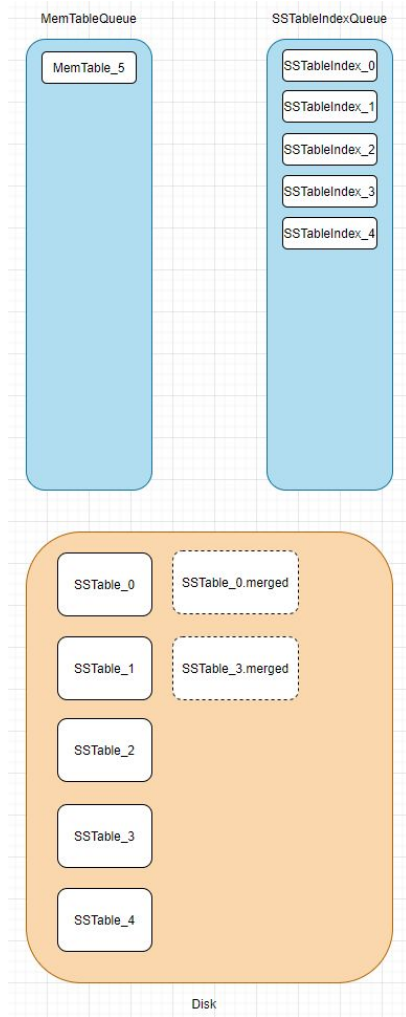


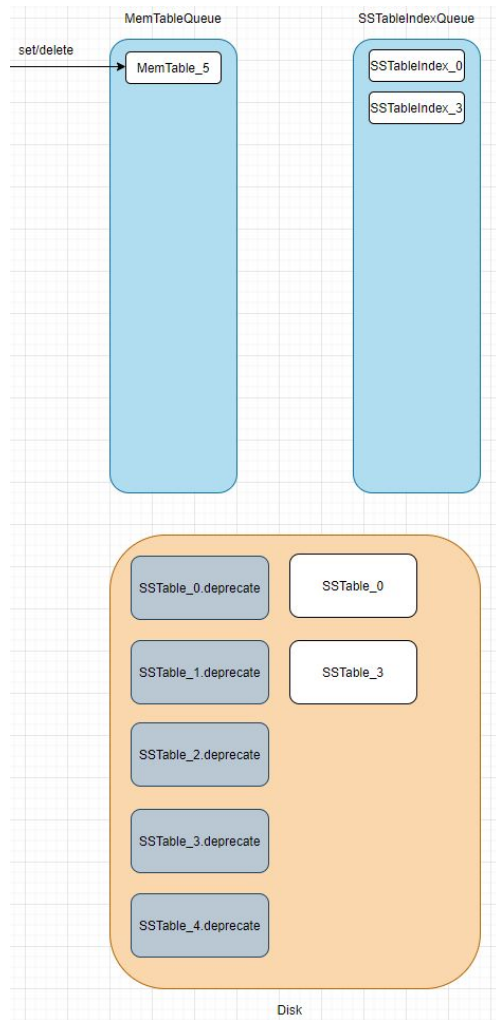


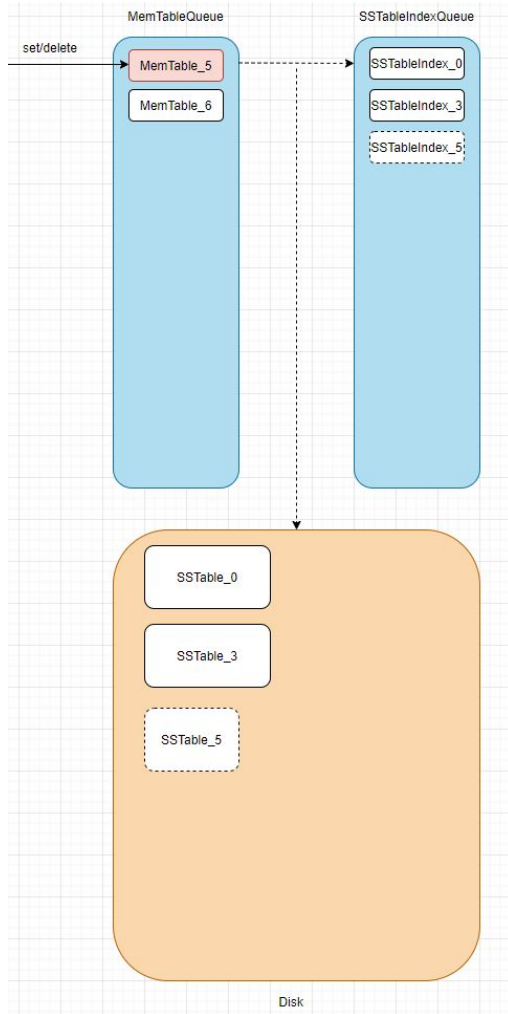












Demo - set and get

```
21 void test() {  
22     db_config::MEMTABLE_APPROXIMATE_MAX_SIZE_IN_BYTES = 10;  
23     db_config::SSTABLE_INDEX_BLOCK_SIZE_IN_BYTES = 8;  
24     db_config::NUM_SSTABLE_TO_COMPACT = 2;  
25     db_config::SSTABLE_APPROXIMATE_MAX_SIZE_IN_BYTES = 20;  
26  
27     ProjectDb db;  
28  
29     for(auto i = 0; i < 12; i++) {  
30         db.set(to_string(i), value: to_string(i) + "-Hello World!");  
31     }  
32  
33     for(auto i = 0; i < 12; i++) {  
34         const auto key = to_string(i);  
35         const auto& val = db.get(key);  
36         if(val.has_value()) {  
37             cout << val.value() << endl;  
38         } else {  
39             cout << "Key: [" << key << "] not found in db!" << endl;  
40         }  
41     }  
42 }
```

```
cd cmake-build-release && ./main  
0-Hello World!  
1-Hello World!  
2-Hello World!  
3-Hello World!  
4-Hello World!  
5-Hello World!  
6-Hello World!  
7-Hello World!  
8-Hello World!  
9-Hello World!  
10-Hello World!  
11-Hello World!
```

Demo - set and get

```
[mli@/home/mli/COMSW_4995_Design_Using_C++/COMSW_4995_Design_Using_Cpp_Project/cmake-build-release/projectdb]  
[$] ls -lh  
total 28K  
-rw-r--r-- 1 mli mli 86 Apr  9 12:32 project_db_1.sst  
-rw-r--r-- 1 mli mli 53 Apr  9 12:32 project_db_11.sst  
-rw-r--r-- 1 mli mli 53 Apr  9 12:32 project_db_12.sst  
-rw-r--r-- 1 mli mli  0 Apr  9 12:32 project_db_12.txt  
-rw-r--r-- 1 mli mli 86 Apr  9 12:32 project_db_3.sst  
-rw-r--r-- 1 mli mli 86 Apr  9 12:32 project_db_5.sst  
-rw-r--r-- 1 mli mli 86 Apr  9 12:32 project_db_7.sst  
-rw-r--r-- 1 mli mli 86 Apr  9 12:32 project_db_9.sst
```

Demo - set, remove and get

```
21 void test() {  
22     db_config::MEMTABLE_APPROXIMATE_MAX_SIZE_IN_BYTES = 10;  
23     db_config::SSTABLE_INDEX_BLOCK_SIZE_IN_BYTES = 8;  
24     db_config::NUM_SSTABLE_TO_COMPACT = 2;  
25     db_config::SSTABLE_APPROXIMATE_MAX_SIZE_IN_BYTES = 20;  
26  
27     ProjectDb db;  
28  
29     for(auto i = 0; i < 12; i++) {  
30         db.set(to_string(i), value: to_string(i) + "-Hello World!");  
31     }  
32  
33     for(auto i = 0; i < 6; i++) {  
34         db.remove(to_string(i));  
35     }  
36  
37     for(auto i = 0; i < 12; i++) {  
38         const auto key = to_string(i);  
39         const auto& val = db.get(key);  
40         if(val.has_value()) {  
41             cout << val.value() << endl;  
42         } else {  
43             cout << "Key: [" << key << "] not found in db!" << endl;  
44         }  
45     }  
46 }
```

```
cd cmake-build-release && ./main  
Key: [0] not found in db!  
Key: [1] not found in db!  
Key: [2] not found in db!  
Key: [3] not found in db!  
Key: [4] not found in db!  
Key: [5] not found in db!  
6-Hello World!  
7-Hello World!  
8-Hello World!  
9-Hello World!  
10-Hello World!  
11-Hello World!
```

Demo - set, remove and get

```
[mli@/home/mli/COMSW_4995_Design_Using_C++/COMSW_4995_Design_Using_Cpp_Project/cmake-build-release/projectdb]  
[$] ls -lh  
total 32K  
-rw-r--r-- 1 mli mli 86 Apr 9 12:29 project_db_1.sst  
-rw-r--r-- 1 mli mli 90 Apr 9 12:29 project_db_11.sst  
-rw-r--r-- 1 mli mli 100 Apr 9 12:29 project_db_13.sst  
-rw-r--r-- 1 mli mli 58 Apr 9 12:29 project_db_15.sst  
-rw-r--r-- 1 mli mli 0 Apr 9 12:29 project_db_15.txt  
-rw-r--r-- 1 mli mli 86 Apr 9 12:29 project_db_3.sst  
-rw-r--r-- 1 mli mli 86 Apr 9 12:29 project_db_5.sst  
-rw-r--r-- 1 mli mli 86 Apr 9 12:29 project_db_7.sst  
-rw-r--r-- 1 mli mli 86 Apr 9 12:29 project_db_9.sst
```

Measurements

- Main features of Database:
 - writing - `.set()`
 - removing - `.remove()`, similar to `.set()`
 - reading - `.get()`
- Supports randomly generated data
- Supports reading data from csv

Data for Benchmark


- Key value pairs stored in a vector
- Randomly generated
 - 0, 1, 2, ... as keys
 - randomly generated string as values
- Trending YouTube Video Statistics
 - video ID as keys: 2kyS6SvSYSE
 - video title as values: "WE WANT TO TALK ABOUT OUR MARRIAGE"

Benchmark Description

- 7 Tasks:
 - fill sequential / random
 - overwrite
 - delete sequential / random
 - read sequential / random
- Total time for task in microseconds / number of entries


Benchmark: Randomly Generated

- First configuration
 - 2 mb in memory
 - 250 kb segments
- Fill faster than overwrite, delete
- Sequential faster than random
- Slower for more entries
- Overwrite?




Benchmark: Read

- Fast for low numbers
- Slow when data size > 2 mb
 - values: 100 characters
 - >2.4 mb for 24,000 entries




How to Solve This?

- 32,000 entries
- Second configuration
 - 2 mb in memory
 - 2 kb segments
- Much faster read
- Similar performance otherwise
- Very small segments?



Benchmark: Trending YouTube Video Statistics

- Using the first configuration
- Smaller value sizes
- Entirely in memory



Further Benchmarks

- Comparison with LevelDB
- Flush to disk
- More configurations

Future Works

- Multiple potential optimizations
 - Implement Bloom filter to achieve better performance for identifying keys that does not exist
 - Implement optimization for sequential read
 - Block can be kept in memory for a while before throwing away
 - It is possible for us to directly perform the merge while flushing to SSTable, instead of merge SSTable in a separate job
- More benchmarks
- Build on top of the library to actually build a distributed key-value store

Questions?

- LSM-tree
- Workflow
- Implementation
- Measurements
- Future Works