

ProjectDb

A Key Value Storage Library

Lieyang Chen
Mengwen Li (ml4643)
Shengtian Mao (sm4954)

<https://github.com/mli9502/ProjectDb>

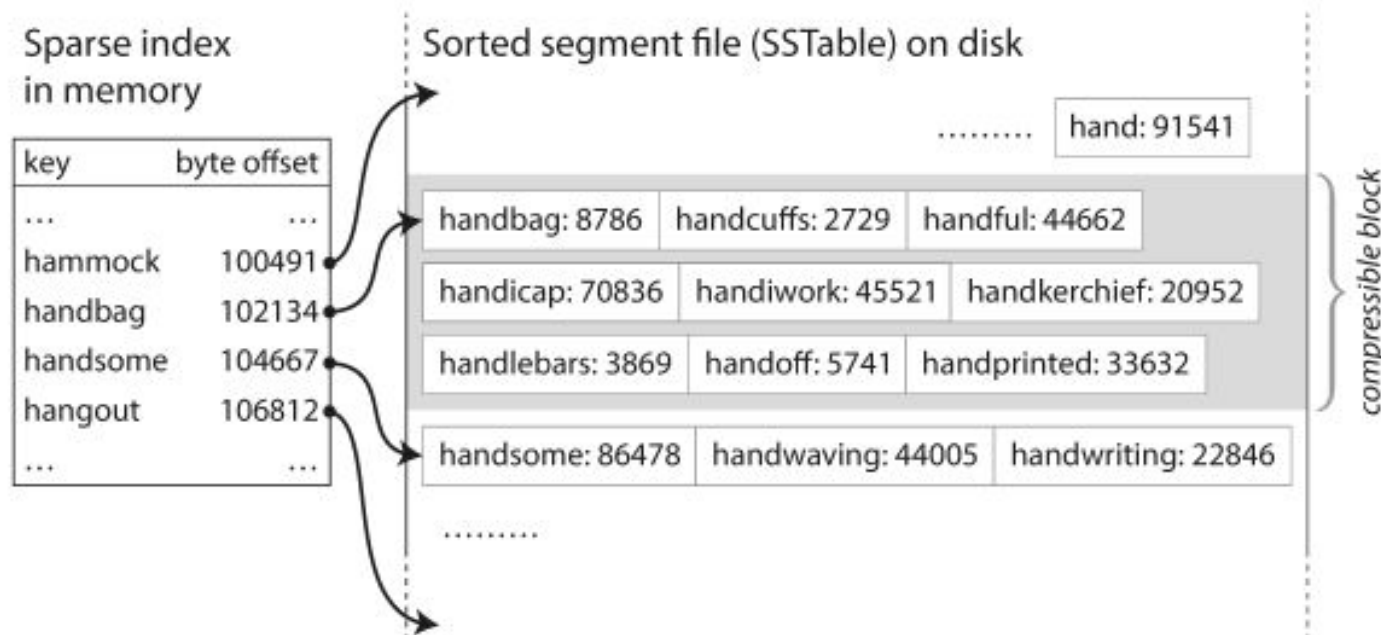
Introduction

- In this project, we implemented a key-value storage engine library by keeping record
- Provides a better performance comparing to directly writing file/reading file from disk.
- Provide a user friendly interface.
-
- Currently Distributed NoSQL database is very popular and widely used, things like Cassandra, MongoDB, ... And a key-value storage engine could be used as a base for these distributed NoSQL databases.

Basic Terminologies

- MemTable: A table containing sorted key-value entries stored in memory.
- SSTable(aka segment file): A table containing sorted key-value entries stored in disk. (Created by flush Memtable into disk).
- Segment: A small-size block of sorted key-value entries stored in disk. Memory reads in a segment instead of entire SSTable when performing read operations. (SSTable is composed of many segments).
- Sparse Index Table: A table that keeps track of the position of the beginning of each Segment.


An intuitive Graph



source: Designing Data-Intensive Applications: p76 - p79

Algorithm

LSM-Tree(Log-structured merge-tree)



source:

https://en.wikipedia.org/wiki/Log-structured_merge-tree#/media/File:LSM_Tree.png

Workflow

- When a write comes in, add it to an in-memory balanced tree data structure (for example, a red-black tree). This in-memory tree is sometimes called a **memtable**.
- When the memtable gets bigger than some threshold—typically a few megabytes—write it out to disk as an **SSTable** file. This can be done efficiently because the tree already maintains the key-value pairs sorted by key. The new SSTable file becomes the most recent segment of the database. While the SSTable is being written out to disk, writes can continue to a new memtable instance.
- In order to serve a read request, first try to find the key in the memtable, then in the most recent on-disk segment, then in the next-older segment, etc.
- From time to time, run a **merging** and compaction process in the background to combine segment files and to discard overwritten or deleted values.

Complexity

Append-only

Deletion is not decreasing file sizes

Complexity

Append-only

Deletion is not decreasing file sizes

Serialization/Deserialization

Append-only

Deletion is not decreasing file sizes

Serialization/Deserialization

Append-only

Deletion is not decreasing file sizes


Serialization/Deserialization


Append-only


Deletion is not decreasing file sizes


Serialization/Deserialization

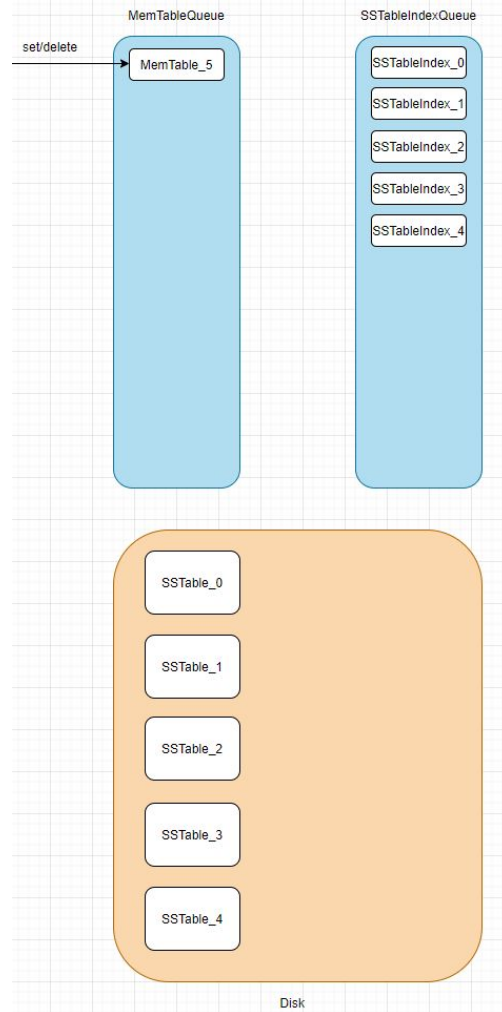
```
71  /**
72   * Serializable follows the following recursive definitions:
73   *
74   * SerializableBase = Trivial |
75   * SerializableBase = SerializableUserDefinedType
76   * Serializable = SerializableBase |
77   * Serializable = Pair<Serializable, Serializable> |
78   * Serializable = Container<Serializable>
79   */
80  // Define type trait for the basic serializable unit.
81  template <typename T>
82  struct serializable_base_trait : std::false_type {};
83
84  template <SerializableBase T>
85  struct serializable_base_trait<T> : std::true_type {};
86  // Define type trait for generic serializable.
87  template <typename T>
88  struct serializable_trait : serializable_base_trait<T> {};
89
90  template <Pair T>
91  struct serializable_trait<T>
92  : conjunction<serializable_trait<remove_const_t<typename T::first_type>>,
93               serializable_trait<typename T::second_type>> {};
94
95  template <Container T>
96  struct serializable_trait<T> : serializable_trait<typename T::value_type> {};
97
98  template <typename T>
99  concept Serializable = serializable_trait<T>::value;
```

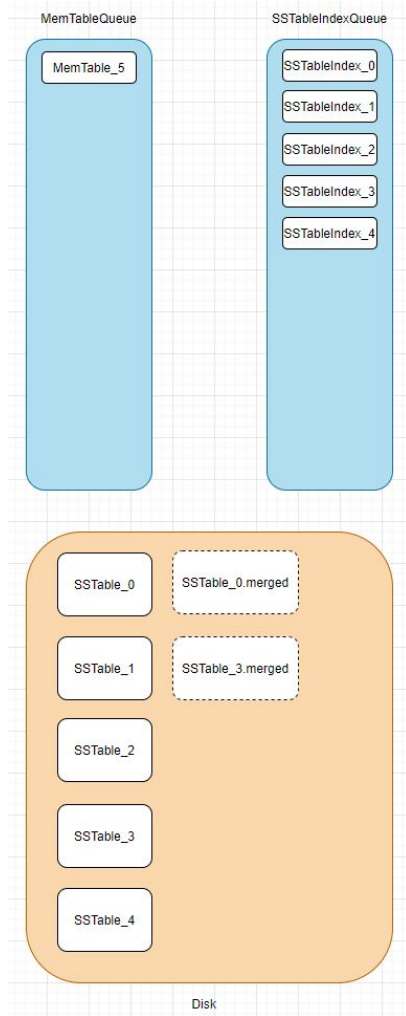


















Measurements

Main features of Database:

- writing - `.set()`

- removing - `.remove()`, which works similarly to `.set()`

- reading - `.get()`

Supports randomly generated data

Supports reading data from csv

Data for Benchmark

Key value pairs stored in a vector

Randomly generated

- 0, 1, 2, ... as keys

- randomly generated string as values

Trending YouTube Video Statistics

- video ID as keys

- video title as values

Benchmark Description

7 Tasks:

fill sequential / random

overwrite

delete sequential / random

read sequential / random

Benchmark: Randomly Generated

First configuration

2 mb in memory

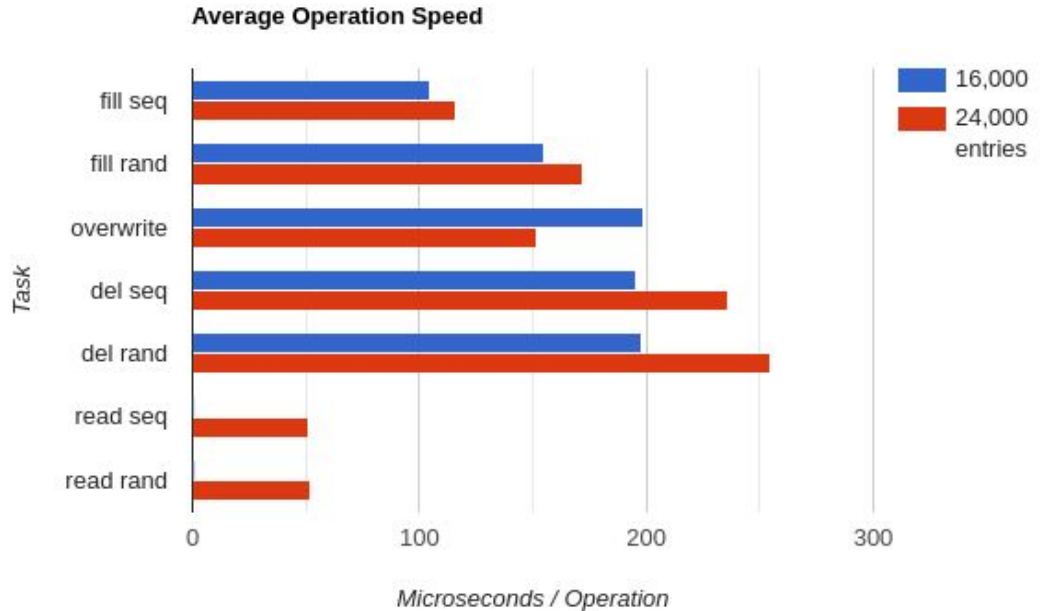
250 kb index block

Slower for more entries

Fill faster than delete

Sequential faster than random

Overwrite?



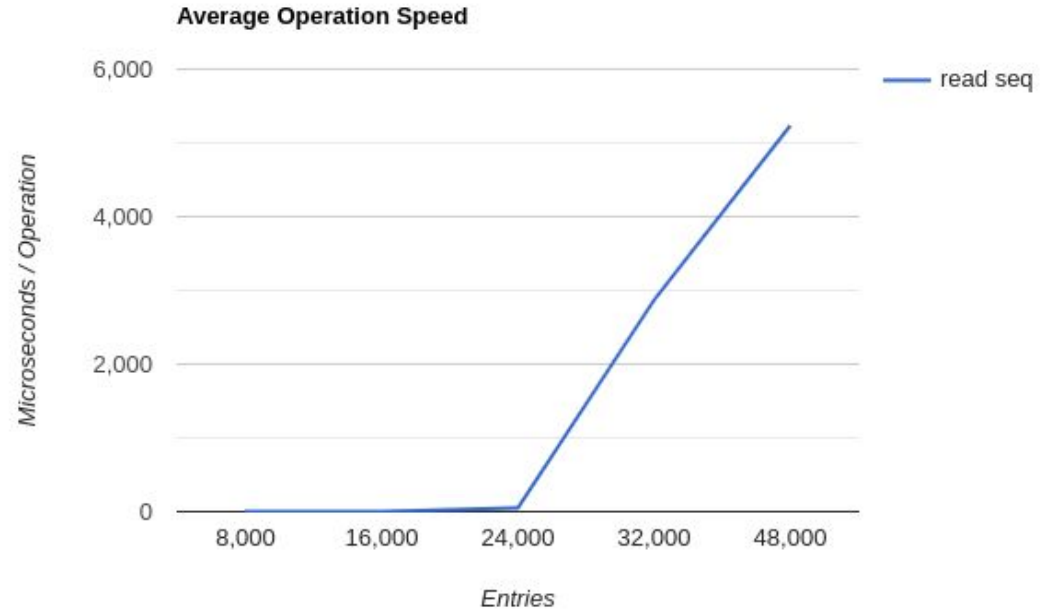
Benchmark: Read

Fast for low numbers

Slow when data size > 2 mb

values: 100 characters

2.4 mb for 24,000 entries



How to Solve This?

32,000 entries

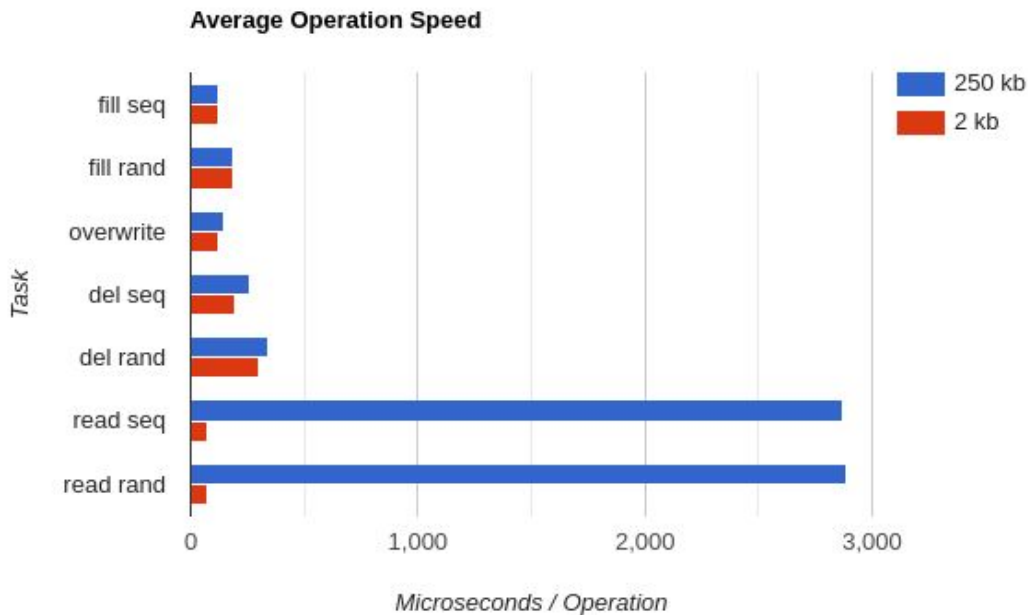
Second configuration

2 mb in memory

2 kb index block

Much faster read

Similar performance otherwise



Benchmark: Trending YouTube Video Statistics


Example:

2kyS6SvSYSE

“WE WANT TO TALK ABOUT OUR MARRIAGE”

Using the first configuration

Smaller value sizes



Further Benchmarks

Comparison with LevelDB

Flush to disk

More configurations

Questions?

- Measurements
 - randomly generated and csv
 - writing, removing, reading
 - 2 kb index block faster